



Deep Learning for Visual Computing

Deep Image Classification (2/2)

Christopher Pramerdorfer
Computer Vision Lab, TU Wien

Deep Image Classification (part 2)

- ▶ Pooling layers
- ▶ Classification backends
- ▶ Network design basics
- ▶ Residual networks

What CNNs learn

Deep Image Classification

Dimensionality Reduction

Recall that conv layers

- ▶ Retain the input size $W_{l-1} \times H_{l-1}$ (padding)
- ▶ Or reduce it only slowly (no padding)

This

- ▶ Slows down computations
- ▶ Leads to more parameters when combined with linear layers

Deep Image Classification

Dimensionality Reduction

CNNs thus include some form of **pooling**

- ▶ Reduce W_{l-1} and H_{l-1} via local aggregation
- ▶ While leaving D_{l-1} unchanged

Some CNNs also do the opposite (e.g. [3])

- ▶ Keep W_{l-1} and H_{l-1} but decrease D_{l-1}
- ▶ Using a conv layer with $c = 1$ and $D_l < D_{l-1}$

Deep Image Classification

Dimensionality Reduction

Most common form is 2×2 max-pooling with stride 2

- ▶ $W_l = W_{l-1}/2$, $H_l = H_{l-1}/2$, and $c = 2$
- ▶ Output of neuron h is $\max(\mathbf{X}_h)$ with $\mathbf{X}_h \in \mathbb{R}^{2 \times 2}$





Image from [cs231n.github.io](https://github.com/cs231n)

Deep Image Classification

Dimensionality Reduction

Number of neurons reduced by factor 4

- ▶ Corresponding efficiency increase
- ▶ At the cost of losing spatial resolution



Deep Image Classification

Dimensionality Reduction

Can also perform **average-pooling**

- Compute $\text{mean}(\mathbf{X}_h)$ instead of $\text{max}(\mathbf{X}_h)$

Or a conv layer with stride $s > 1$




Image from github.com

Deep Image Classification

Classification Backends

Conv and pooling layers produce 3D tensors ($W_l \times H_l \times D_l$)

For classification we convert to vectors \mathbf{x}_l

- ▶ Option 1: flatten tensor like we did with images
- ▶ Option 2: **global average-pooling** with size $W_l \times H_l$

Deep Image Classification

Classification Backends

Allows us to connect linear layers, resulting in

- ▶ A linear or non-linear (MLP) classifier
- ▶ That processes vectors of **learned** features

Modern architectures often use a linear classifier

- ▶ So the learned features are so powerful
- ▶ That the simplest classifier is sufficient

We finally have task-specific high-level features!

Deep Image Classification

Network Architecture Design

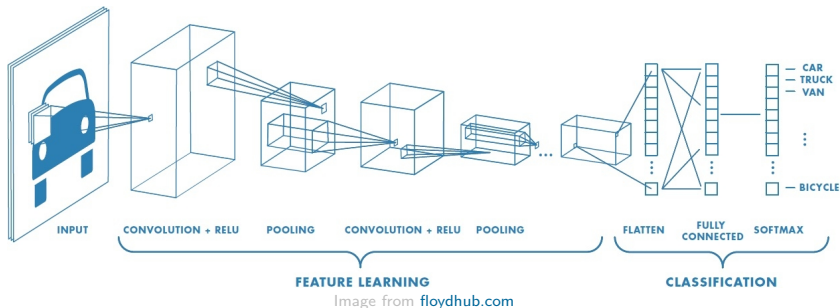
This concludes the basic layer types and purposes

- ▶ Conv + ReLU layers for feature extraction
- ▶ Pooling layers for dimensionality reduction
- ▶ Linear layers for classification

Deep Image Classification

Network Architecture Design

The question is how to arrange these layers properly



Deep Image Classification

Network Architecture Design

This is a complex topic

- ▶ Most layers have (several) hyperparameters
- ▶ Layer arrangement is also a hyperparameter

Cannot just search for good hyperparameters

- ▶ Training once can take hours or even days

But tried and true practices exist

- ▶ Next slides will introduce a basic design recipe

Deep Image Classification

Network Architecture Design – Basic Recipe

Use square images, $H_0 = W_0 = R$

- ▶ Resize images such that smaller side has size R
- ▶ Then extract a center crop of size $R \times R$

R should be divisible by 2 many times

- ▶ Avoid problems during pooling

Deep Image Classification

Network Architecture Design – Basic Recipe

Start with small R

- ▶ And test if increasing R makes sense
- ▶ $R = 224$ is popular for classification
- ▶ But much smaller R might be sufficient

Get R below 100 quickly via aggressive pooling

- ▶ To improve efficiency
- ▶ $R = 224$: conv with $c = 7, s = 2 \Rightarrow 2 \times 2$ max-pooling

Deep Image Classification

Network Architecture Design – Basic Recipe

Use conv \Rightarrow conv \Rightarrow pooling blocks

- ▶ Conv layers with $c = 3$, stride 1, padding, and ReLUs
- ▶ Pooling layers with 2×2 max-pooling with stride 2

Start with 32 or 64 feature maps

- ▶ Increase by factor 2 in each subsequent block
- ▶ Up to a maximum of 512 feature maps

Stack blocks until $R \leq 7$

- ▶ This means 4 or 5 such blocks

Deep Image Classification

Network Architecture Design – Basic Recipe

Finish with linear classifier

- ▶ Global average pooling
- ▶ Followed by linear layer with T neurons

This recipe is a good starting point


- ▶ Decent performance on many datasets (try yourself!)
- ▶ Tune some hyperparameters depending on time budget

Deep Image Classification

Network Architecture Design – Basic Recipe

Example result with $R = 224$

- ▶ Not suited for our cat vs. dog problem as $R = 32$
- ▶ Design and compare suitable architectures in Assignment 2



Deep Image Classification

Network Architecture Design

Example network has a depth of 10 (layers with parameters)

- ▶ Should we go deeper?



Image from memegen.com

Deep Image Classification

Network Architecture Design

In theory we can go as deep as we want

- ▶ Simply by stacking conv layers

In practice as a rule of thumb

- ▶ Going deeper improves performance
- ▶ With proper regularization (more later)
- ▶ Up to a depth of around 15 with basic network designs

Deep Image Classification

Network Architecture Design

Until recently deeper architectures were uncommon

- ▶ Hard to train due to **vanishing/exploding gradients**
- ▶ Caused by aggregating small/large gradients

Recall that gradient computation entails chained multiplications

- ▶ If the local gradients are small (or large)
- ▶ The final gradient may be close to 0 (or very large)

Deep Image Classification

Network Architecture Design

This problem has been largely resolved via

- ▶ ReLU activation functions
- ▶ Proper parameter initialization
- ▶ Intermediate normalization via Batch Normalization

In order to preserve the signal strength throughout the net

- ▶ By making layers preserve input variance

Deep Image Classification

Network Architecture Design

The input variance is

- ▶ The variance of neuron inputs (over the training dataset)
- ▶ Usually over whole layer or per feature map (conv)

Depends on input (image) preprocessing too

- ▶ E.g. normalizing from $[0, 255]$ to $[-1, 1]$ reduces variance

So make sure to **normalize** input images

- ▶ Subtract per-channel mean of training set
- ▶ Then divide by per-channel standard deviation

Deep Image Classification

Parameter Initialization

To summarize

- ▶ Bias values are uncritical and can be set to 0
- ▶ Weights should be sampled from a normal distribution
- ▶ With $\mu = 0$ and σ set to preserve the input variance

Best method with ReLU activations is [2]

- ▶ PyTorch: `nn.init.kaiming_normal_`

Deep Image Classification

Batch Normalization

Proper initialization preserves input variance only initially

- ▶ Parameters change during training
- ▶ Thus output distribution of layer changes over time

This complicates training

- ▶ Must account for changes in input distribution
- ▶ Layer input affected by parameters of all previous layers

Deep Image Classification

Batch Normalization

Batch normalization reduces this problem

- ▶ Estimate input mean and variance
- ▶ Using current minibatch to approximate training set
- ▶ Normalize accordingly (per feature map for conv layers)

Should be done before activation function

- ▶ PyTorch: add layer between conv and ReLU layers

Deep Image Classification

Batch Normalization

Advantages of batch normalization

- ▶ Improves robustness to bad initialization
- ▶ Permits higher learning rates (e.g. 0.1)
- ▶ Has a regularizing effect (more later)

Suggestions

- ▶ Use after all conv and hidden linear layers
- ▶ Shuffle training set before every epoch

Deep Image Classification

Residual Networks

With gradient problems out of our way, how deep should we go?



Image from knowyourmeme.com

Deep Image Classification

Residual Networks

Going deeper should not harm *training* performance

- ▶ Increases the model capacity
- ▶ Network can learn identity mappings to “skip” layers

However in practice this is the case for depths > 30 or so

- ▶ Very deep nets become hard to train
- ▶ This is not due to vanishing gradients

Residual Networks (ResNets) solve this problem

- ▶ Current state of the art architecture

Deep Image Classification

Residual Networks

ResNets consist of multiple **residual blocks**

- Below is the original version, others exist





Image adapted from [1]

Deep Image Classification

Residual Networks

Learn additive residual function \mathcal{F} with respect to x

- ▶ Provides guidance (learn what to add/remove from x)
- ▶ ResNets support depths of 1000 layers and more

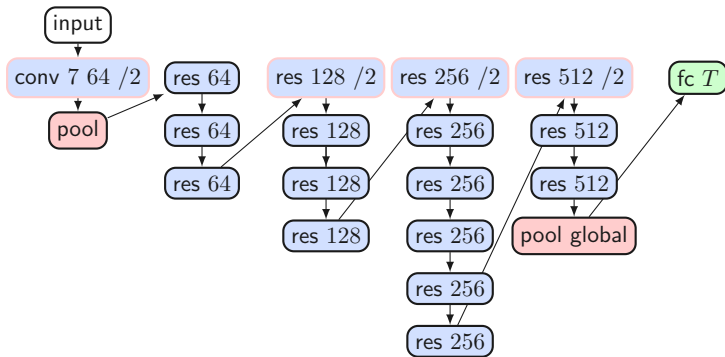


Deep Image Classification

Residual Networks

ResNet-34 for ImageNet ($R = 224$)

- First and last layers as already covered



What CNNs Learn

CNN classifiers are trained like other NN classifiers

- ▶ Softmax and cross-entropy loss
- ▶ Gradient descent and backpropagation

CNN classifiers thus *simultaneously* learn

- ▶ How to extract good features (conv layer parameters)
- ▶ How to perform (linear) classification using these features
- ▶ This joint optimization is why CNNs are so powerful

What CNNs Learn

First conv layer usually learns Gabor-like filters

- ▶ Similarly to early human vision (!)



Image from [cs231n.github.io](https://github.com/cs231n)

What CNNs Learn

Later conv layers learn to respond to more specific concepts

- ▶ The high-level features we want




Image adapted from [4]

Bibliography

- [1] *Deep Residual Learning for Image Recognition*. CVPR. 2016.
- [2] Kaiming He et al. *Delving deep into rectifiers: Surpassing human-level performance on imagenet classification*. CVPR. 2015, pp. 1026–1034.
- [3] Christian Szegedy et al. *Going Deeper with Convolutions*. CVPR. 2015.
- [4] *Visualizing and understanding convolutional networks*. ECCV (2014).